Native E0: physical KV address sharing

Same Q, token order, plan, scratch and output buffers; only installed KV page IDs differ.

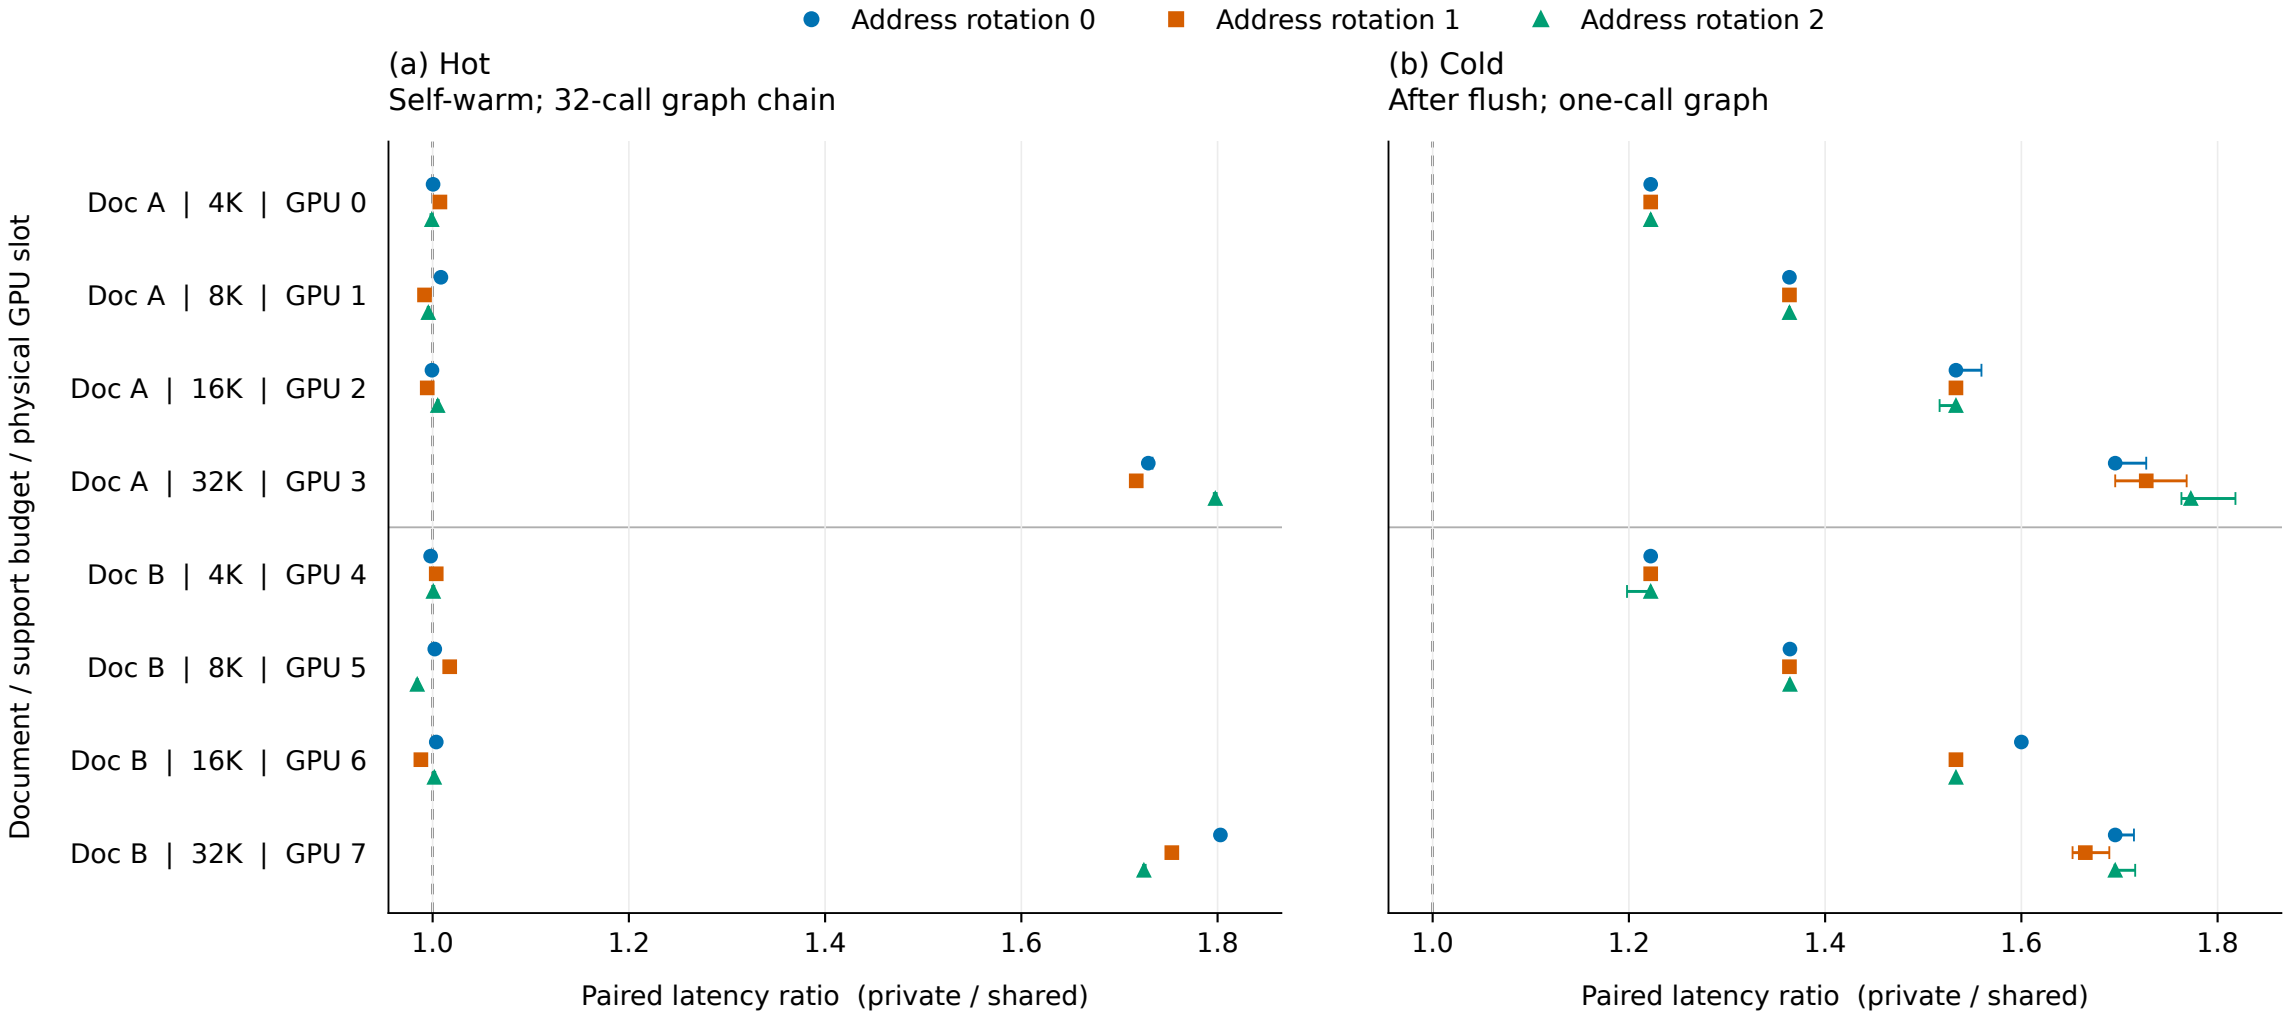

Points: median private/shared ratio; whiskers: 95% bootstrap CI over 30 paired blocks (10,000 draws).

Random/reverse measurements stay clustered. Ratio > 1 means the shared-address arm is faster.

Each doc × budget cell ran on a different GPU: inference is within-cell only; budgets are not connected.

Cold timings show 2.048 µs quantization. Intervals omit systematic/quantization uncertainty; some collapse.

BF16 KV, B3 / GQA8 / Hkv2, layer 15, step 64; O + output copy timed; index installation excluded.

Doc A: docfinqa256; Doc B: publicdev1_256. This is not an application gain or a measurement of L2/DRAM traffic.